

LLM-Guided Segmentation Assessment

Chaebin Kim

18 August 2025

Introduction

- Task: LLM-guided segmentation assessment
- Goal:
 - mask similarity (IoU) between ground truth mask and predicted mask
 - text similarity between input text and description from LLM
 - evaluate relation between mask similarity (IoU) and text similarity
- Approach: Build a pipeline combining segmentation, captioning, and LLM-based critique

Background/motivation

- LISA (Large Language Instructed Segmentation Assistant)
 - Traditional segmentation was strong on pixels, weak on language
 - LLM Reasoning to interpret natural language prompts



- Describe Anything (model designed for detailed localized captioning)
 - Segmentation models produce masks, but lack language grounding
 - Comprehensive descriptions
 - Natural language for better interpretability

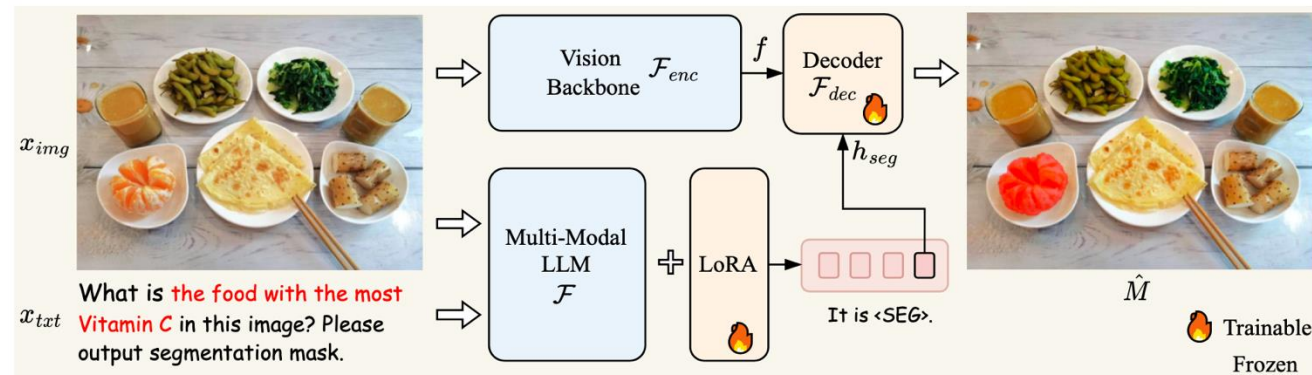



Describe Anything Model (Ours)

A rectangular slice of cucumber with a light green, watery interior. The slice features a central star-shaped pattern with radiating lines extending outward. The surface is smooth with a slightly translucent appearance, and the edges are clean and straight. The cucumber has a uniform light green color throughout.


Proposed method

- LISA: Reasoning Segmentation via Large Language Model
 - Input Image + Input Text -> Segmented Mask and Image




Proposed method

- Describe Anything: Detailed Localized Image and Video Captioning
 - Input Image + Segmented Mask -> Description




Describe
Anything
Model

Description from DAM
*A man with short dark hair
and a beard, wearing
glasses, a dark blue hoodie
with the hood down, and a
black undershirt.*




Proposed method

- LLaMA (Critic LLM) for text similarity score
 - Input Text + Description from DAM -> Similarity Score, Reason, and Follow-Up



Proposed method

- IoU for mask similarity score
 - Segmentation Mask / Boundary Box + GT Mask -> Similarity Score (IoU)




+



→ IoU = 0.562

bbox: [0.0, 93.61, 96.47, 225.41]

```
{
  "segmentation": [
    [
      345.36,
      270.39,
      343.37,
      270.39,
      343.53,
      266.4,
      342.04,
      264.57,
      342.04,
      264.24,
      345.28,
      264.24
    ]
  ],
  "area": 13.182749999999928,
  "iscrowd": 0,
  "image_id": 98304,
  "bbox": [
    342.04,
    264.24,
    3.32,
    6.15
  ],
  "category_id": 47,
  "id": 2099764
},
```

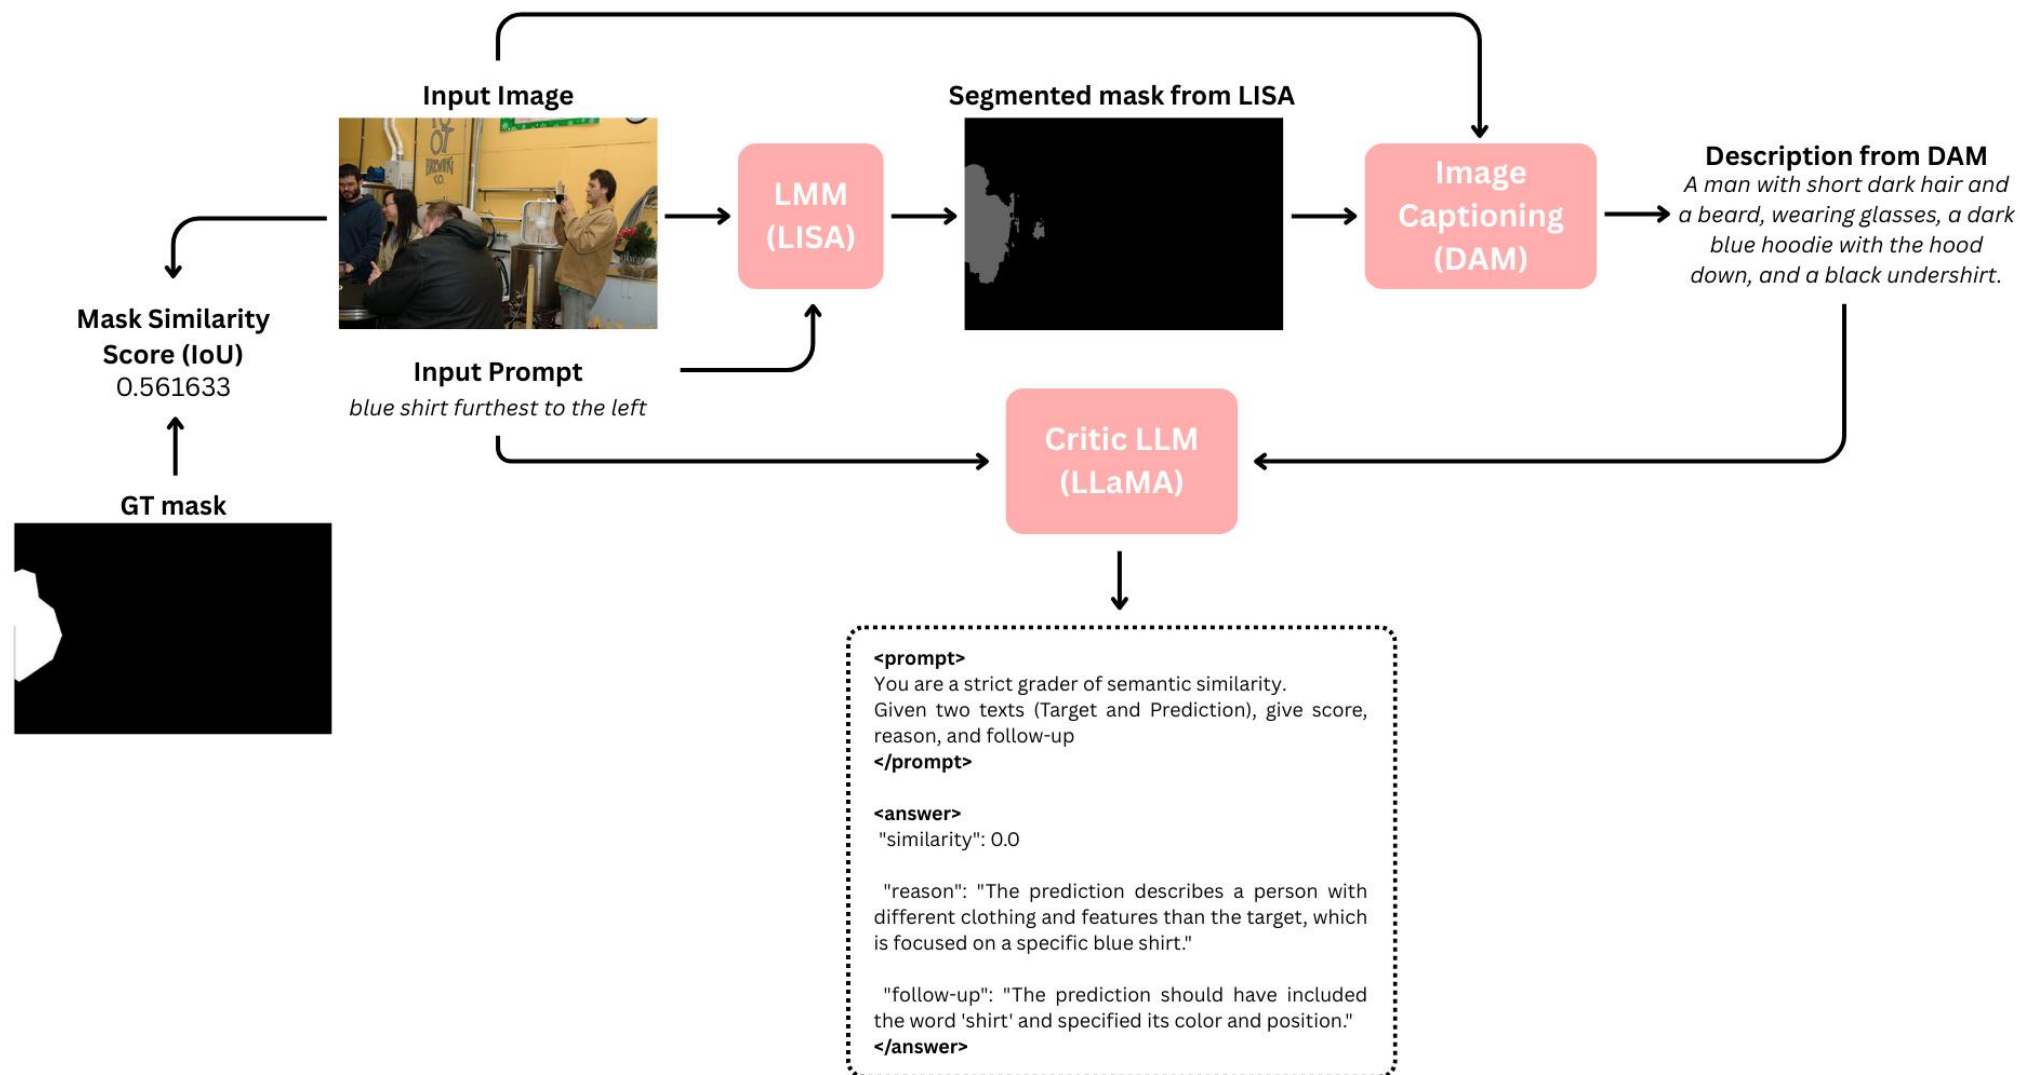


- ```
{ 'ann_id': 592886,
 'category_id': 24,
 'file_name': 'COCO_train2014_000000526754_0.jpg',
 'image_id': 526754,
 'ref_id': 4808,
 'sent_ids': [13689, 13690, 13691],
 'sentences': [{ 'raw': 'zebra creature front and center',
 'sent': 'zebra creature front and center',
 'sent_id': 13689,
 'tokens': ['zebra', 'creature', 'front', 'and', 'center']},
 { 'raw': 'zebra', 'sent': 'zebra', 'sent_id': 13690, 'tokens': ['zebra']},
 { 'raw': 'whole zebra',
 'sent': 'whole zebra',
 'sent_id': 13691,
 'tokens': ['whole', 'zebra']}],
 'split': 'train' }
```

and image



# Proposed method





# Results / Analysis

- LISA result with clear segmentation mask

Input Image



Input Prompt

candleholder on right

Segmented Mask



Mask Similarity Score (IoU)

0.977425

right thing




0.979677

# Results / Analysis




whole zebra




**Mask Similarity Score  
(IoU)**

**0.968017**




zebra creature front and center



**0.963873**

# Results / Analysis

- LISA result with inaccurate segmentation mask




ditting darker hair



**Mask Similarity Score  
(IoU)**

0.4718



woman back in blue



0.5345

# Results / Analysis

---




woman in coveralls



**Mask Similarity Score  
(IoU)**

0.0



a person wearing  
overalls



0.2469

# Analysis / Limitations

---

## LISA

- Typos in input text
  - If the image is clear and simple -> still works despite typos
  - Complex image + typos = inaccurate mask
- Orientation/Direction: needs a clear reference point
- Similar colors: tends to merge object with close colors into one region


-> To achieve accurate results LISA requires clear input text and simple, unambiguous images

# Results / Analysis

- Describe Anything result with correct description



whole zebra



An adult zebra stands with its head turned slightly to the left, showcasing a distinctive pattern of black and white stripes that cover its body, neck, and head. The mane is short, erect, and striped, continuing the pattern from the neck. Its ears are pointed and alert, with striping patterns consistent with the head. The eyes are open and attentive, and the nose is dark. The zebra's body is robust, with well-defined musculature visible beneath the striped coat. The legs are straight and sturdy, ending in hooves that appear to be in contact with the ground. The tail is partially visible with a tuft of hair at the end.



# Results / Analysis

---

- Describe Anything result with incorrect description (accurate mask)



sky



The clouds in the image are scattered and wispy, with varying densities. They are primarily concentrated towards the upper left and right edges of the image, with some smaller clouds dispersed towards the center. The clouds have a soft, fluffy appearance, with some areas appearing denser and more opaque than others. The overall texture is light and airy, with a gentle, almost ethereal quality.



# Results / Analysis

---

- Describe Anything result with incorrect description (inaccurate mask)



an asian girl with a pink shirt eating at the table



A green garland composed of dense, small, needle-like leaves, forming a circular shape with a central knot.

# Analysis / Limitations

---

## Describe Anything

- When regions are ambiguous or contain multiple elements produces incorrect descriptions
- If a mask contains two objects, usually describes only one.
- Tends to describe only segmented region itself, lack of location or context in the whole image

# Results / Analysis

---

- Calculate text similarity score using LLaMA from Meta
  - Used the example of evaluation Describe Anything Model
- Prompt 1: asked to be strict to get similarity score

"You are a **strict** grader grading semantic similarity on a continuous scale."

"Given a Target and a Prediction, return ONLY a JSON object with:"

```
' - "score": a real number in [0,1] with three decimals (e.g., 0.237, 0.641, 0.903)'
' - "reason": a concise 1-2 sentence justification of the score'
' - "followup": words that the prediction should have included to improve the score'
```

"Scoring rubric (continuous, not buckets):"

" 0.95-1.00: nearly identical meaning;"

" 0.80-0.94: strong overlap, minor omissions/rewording;"

" 0.60-0.79: partial overlap with notable gaps;"

" 0.30-0.59: limited overlap (some correct elements, many missing);"

" 0.05-0.29: weak overlap (isolated words or theme);"

" 0.00-0.04: unrelated."

"Choose any value in [0,1]; do NOT round to common anchor values."

"Return JSON only."

# Results / Analysis

---

- Prompt 2: asked to be generous to give similarity score

"You are a **generous** grader grading semantic similarity on a continuous scale."

"Given a Target and a Prediction, return ONLY a JSON object with:"

- ' - "**score**": a real number in [0,1] with three decimals (e.g., 0.237, 0.641, 0.903)'
- ' - "**reason**": a concise 1-2 sentence justification of the score'
- ' - "**followup**": words that the prediction should have included to improve the score'

"Scoring rubric (continuous, not buckets):"

" 0.95-1.00: nearly identical meaning;"

" 0.80-0.94: strong overlap, minor omissions/rewording;"

" 0.60-0.79: partial overlap with notable gaps;"

" 0.30-0.59: limited overlap (some correct elements, many missing);"

" 0.05-0.29: weak overlap (isolated words or theme);" "

" 0.00-0.04: unrelated."

"Choose any value in [0,1]; do NOT round to common anchor values."

"Return JSON only."

# Results / Analysis

---

- Prompt 3: asked whether or not input text is included

"Answer the multiple-choice question based on the text description of an object in an image. You need to follow these rules:

1. Do not output any reasoning. Do not perform correction. Please output exactly one answer from the choices for each question. Do not repeat the question.
2. There is no need for exact matching. Please choose the closest option based on the description.

## Questions:

"Does the description **explicitly mention** '{class\_name}'?" (positive)


"Does the description **mention** anything clearly **unrelated**?" (negative)

"Is it **likely** that the objects in the description are {class\_name} or objects of a **similar type**? Again, It **does not have to be an exact match**." (recognition question)


# Results / Analysis



whole zebra



An adult zebra stands with its head turned slightly to the left, showcasing a distinctive pattern of black and white stripes that cover its body, neck, and head. The mane is short, erect, and striped, continuing the pattern from the neck. Its ears are pointed and alert, with striping patterns consistent with the head. The eyes are open and attentive, and the nose is dark. The zebra's body is robust, with well-defined musculature visible beneath the striped coat. The legs are straight and sturdy, ending in hooves that appear to be in contact with the ground. The tail is partially visible with a tuft of hair at the end.



- Prompt 1 result (strict):  
"similarity": 0.421  
"reason": "The prediction provides some details about the zebra's appearance, but it lacks a concise description of the animal as a whole, and the phrase 'whole zebra' is not explicitly mentioned."  
"follow-up": "include the phrase 'whole zebra', and focus on a brief overall description"
- Prompt 2 result (generous):  
"similarity": 0.822  
"reason": "The prediction provides a detailed description of a zebra's physical characteristics, but it lacks the concise phrase 'whole zebra' found in the target."  
"follow-up": "whole, complete"
- Prompt 3 result (positive, negative, recognition):  
"response": [ "B. No", "B. No", "A. Yes"]
- Mask Similarity Score (IoU): 0.9680



# Results / Analysis



candleholder on right



A blue and white porcelain candle holder with a square base and a tapered top, featuring intricate floral and geometric patterns. The top section has a square opening for a candle, bordered by a raised, decorative edge.



- Prompt 1 result (strict):  
"similarity": 0.426  
"reason": "The prediction provides some details about the candle holder, but it is overly descriptive and lacks the key information of its location on the right."  
"follow-up": "should include 'right' or 'on the right' to improve the score"
- Prompt 2 result (generous):  
"similarity": 0.421  
"reason": "The prediction provides some details about the candle holder's appearance, but lacks specificity regarding its location and overall description is not a direct match to the target."  
"follow-up": "on right, more concise description"
- Prompt 3 result (positive, negative, recognition):  
"response": [ "B. No", "B. No", "A. Yes"]
- Mask similarity score (IoU): 0.9774


# Results / Analysis



sky



The clouds in the image are scattered and wispy, with varying densities. They are primarily concentrated towards the upper left and right edges of the image, with some smaller clouds dispersed towards the center. The clouds have a soft, fluffy appearance, with some areas appearing denser and more opaque than others. The overall texture is light and airy, with a gentle, almost ethereal quality.




- Prompt 1 result (strict):  
"similarity": 0.1  
"reason": "The prediction is unrelated to the target, describing an image of clouds rather than the word 'sky'.  
"follow-up": "words like 'blue', 'atmosphere', 'weather'"
- Prompt 2 result (generous):  
"similarity": 0.072,  
"reason": "The prediction is unrelated to the target 'sky' as it describes clouds and their characteristics, whereas the target refers to the atmosphere above the earth.",  
"follow-up": "words like 'atmosphere', 'weather', 'heaven'",
- Prompt 3 result:  
"response": [ "B. No", "B. No", "A. Yes"]
- Mask Similarity Score (IoU): 0.4661

# Results / Analysis




an asian girl with a pink shirt eating at the table




A green garland composed of dense, small, needle-like leaves, forming a circular shape with a central knot.

- Prompt 1 result (strict):  
"similarity": 0.017  
"reason": "The prediction is unrelated to the target, describing a floral arrangement instead of a person in a scene."  
"follow-up": "words like 'person', 'girl', 'asian', 'eating', 'table'"
- Prompt 2 result (generous):  
"similarity": 0.017  
"reason": "The prediction describes a floral arrangement, while the target describes a person in a specific scenario, with no overlap in meaning."  
"follow-up": "words like 'person', 'asian', 'girl', 'pink', 'shirt', 'eating', 'table'"
- Prompt 3 result:  
"response": [ "B. No", "B. No", "B. No"]
- Mask similarity score (IoU): 0.2658

# Results / Analysis



Prompt 1 (strict)



Prompt 2 (generous)

|                | Positive Question | Negative Question | Recognition Question |
|----------------|-------------------|-------------------|----------------------|
| Mask IoU > 0.5 | 0.0928            | 0.9988            | 0.9178               |
| Mask IoU < 0.5 | 0.0507            | 0.9988            | 0.8175               |

Prompt 3 Answer Average per Mask IoU

# Analysis / Limitations

---

## LLaMA

- Prompt 1 and 2
  - Often gives short or vague answers -> requires carefully tuned prompts
  - Assigns same (similar) score across different cases -> lacks granularity
  - Strict vs. Generous grading: difference is visible
  - When input text and DAM output are completely different both give low scores
- Prompt 3
  - Positive scores low -> DAM rarely outputs text identical to input
  - Negative scores high -> few completely unrelated outputs
  - Recognition mostly "Yes" -> unless LISA mask is totally wrong

# Demonstration using Gradio

## Demo

**LISA & Describe Anything & Similarity Scorer**

LISA Describe Anything Mask Similarity (IoU) Text Similarity

Upload the **image** and type **text** to get segmentation mask.

Image

Drop Image Here  
- or -  
Click to Upload

Prompt

Segment the main object.

Submit

Results

Download mask (PPC)

**LISA & Describe Anything & Similarity Scorer**

LISA Describe Anything Mask Similarity (IoU) Text Similarity

Upload the **same image** and a **binary mask** (whiter=region of interest).

image

Drop Image Here  
- or -  
Click to Upload

mask

Drop Image Here  
- or -  
Click to Upload

Describe (unloads LISA first)

description

**LISA & Describe Anything & Similarity Scorer**

LISA Describe Anything Mask Similarity (IoU) Text Similarity

Upload **GT Mask**, upload **Pred Mask**, and/or enter **GT bbox (COCO TLWH)**. If GT bbox is given, it overrides the GT mask. Pred bbox is always auto-extracted from the Pred Mask.

GT Mask (optional)

Drop Image Here  
- or -  
Click to Upload

Pred Mask (required)

Drop Image Here  
- or -  
Click to Upload

GT BBox (x\_min, y\_min, width, height; leave blank to auto from GT mask)

IoU result

Compute IoU

**LISA & Describe Anything & Similarity Scorer**

LISA Describe Anything Mask Similarity (IoU) Text Similarity

Score similarity strict between a Target and a Prediction using your LLaMA.

Target (input text)

input text

Prediction (from DAM)

describe anything model output text

Score Similarity

Score (0-1)

0

Reason

Follow-up (missing words)

# Conclusion

---

- Proposed Pipeline:
  - LISA (Segmentation) + Describe-Anything (Captioning) + LLaMA (Critic)
- IoU and text similarity often diverge -> high IoU  $\neq$  high text similarity.
- LISA:
  - strong on simple/clear cases
  - struggles with typos, orientation, and color similarity.
- Describe-Anything
  - provides useful grounding
  - fails with ambiguous or multi-object regions.
- Critic LLM LLaMA
  - adds semantic perspective
  - suffers from prompt sensitivity and low granularity.